



香港中文大學 (深圳)  
The Chinese University of Hong Kong

# CSC3100 Data Structures

## Lecture 14: Binary search tree

Yixiang Fang  
School of Data Science (SDS)  
The Chinese University of Hong Kong, Shenzhen

---



# Outline




---

- ▶ In this lecture, we will learn
  - Binary search tree (BST)
  - Operations on BST
    - Search a key, find the minimum/maximum, find successor/predecessor
    - Insert, delete
  - Exercises



# Binary search tree (BST) property

- ▶ BST is a tree such that for each node  $T$ ,
  - the key values in its left subtree are **smaller** than the key value of  $T$
  - the key values in its right subtree are **larger** than the key value of  $T$





# Applications of BST

---

- ▶ Many applications due to its **ordered structure**
  - Useful for indexing and multi-level indexing
  - Helpful in maintaining a sorted stream of data
  - Helpful to implement various searching algorithms and data structures (e.g., TreeMap, TreeSet, Priority queue)

java.util

**Class TreeMap<K,V>**

java.lang.Object

java.util.AbstractMap<K,V>

java.util.TreeMap<K,V>

java.util

**Class TreeSet<E>**

java.lang.Object

java.util.AbstractCollection<E>

java.util.AbstractSet<E>

java.util.TreeSet<E>



# BST ADT

---


- ▶ Support many dynamic set operations
  - searchKey, findMin, findMax, predecessor, successor, insert, delete
- ▶ Running time of basic operations on BST
  - On average:  $\Theta(\log n)$ 
    - The expected height of the tree is  $\log n$
  - In the worst case:  $\Theta(n)$ 
    - The tree is a linear chain of  $n$  nodes



# Searching for a key

- ▶ Given a pointer to the root of a tree and a key  $k$ :
  - Return a pointer to a node with key  $k$  if one exists, otherwise return NIL

- ▶ Example



- ▶ Search for key 13:
  - $15 \rightarrow 6 \rightarrow 7 \rightarrow 13$




# Searching for a key

---

find(x, k)

1. if  $x = \text{NIL}$  or  $k = \text{key}[x]$
2.     **return**  $x$
3. if  $k < \text{key}[x]$
4.     **return** find(left [x], k )
5. **else**
6.     **return** find(right [x], k )



Running time:  $O(h)$ , where  $h$  is the height of tree




# Finding the minimum

- ▶ Goal: find the minimum value in a BST
  - Following **left child pointers** from the root, until a NIL is encountered

findMin(x)

1. **while** left [x]  $\neq$  NIL
2.     **do** x  $\leftarrow$  left [x]
3. **return** x



Minimum = 2

Running time:  $O(h)$ , where  $h$  is the height of tree






# Finding the maximum

- ▶ Goal: find the maximum value in a BST
  - Following **right child pointers** from the root, until a NIL is encountered

findMax(x)

1. **while** right [x]  $\neq$  NIL
2.     **do** x  $\leftarrow$  right [x]
3. **return** x



Maximum = 20


Running time:  $O(h)$ , where  $h$  is the height of tree



# Successor

**Def:**  $\text{successor}(x) = y$ , such that  $\text{key}[y]$  is the smallest key  $> \text{key}[x]$

- ▶ **E.g.:**  $\text{successor}(15) = 17$   
 $\text{successor}(13) = 15$   
 $\text{successor}(9) = 13$




- ▶ **Case 1: right (x) is non empty**
  - $\text{successor}(x) = \text{the minimum in right}(x)$
- ▶ **Case 2: right (x) is empty**
  - go up the tree until the current node is a left child:  $\text{successor}(x)$  is the parent of the current node
  - if you cannot go further (and you reached the root):  $x$  is the largest element



# Successor

successor( $x$ )

1. if right [ $x$ ]  $\neq$  NIL
2.     return findMin(right [ $x$ ])
3.  $y \leftarrow p[x]$
4. while  $y \neq$  NIL and  $x =$  right [ $y$ ]
5.     do  $x \leftarrow y$
6.      $y \leftarrow p[y]$
7. return  $y$




Running time:  $O(h)$ , where  $h$  is the height of tree



# Predecessor

**Def:** predecessor ( $x$ ) =  $y$ , such that key [ $y$ ] is the biggest key  $<$  key [ $x$ ]

- ▶ **E.g.:** predecessor (15) = 13  
predecessor (9) = 7  
predecessor (7) = 6



- ▶ **Case 1: left ( $x$ ) is non empty**
  - predecessor ( $x$ ) = the maximum in left ( $x$ )
- ▶ **Case 2: left ( $x$ ) is empty**
  - go up the tree until the current node is a right child: predecessor ( $x$ ) is the parent of the current node
  - if you cannot go further (and you reached the root):  $x$  is the smallest element



# Insertion





---

- ▶ Goal: Insert value  $v$  into a binary search tree
- ▶ Find the position and insert as a leaf:
  - If  $\text{key}[x] < v$  move to the right child of  $x$ ,  
else move to the left child of  $x$
  - When  $x$  is NIL, we found the correct position
  - If  $v < \text{key}[y]$  insert the new node as  $y$ 's left child  
else insert it as  $y$ 's right child
- Beginning at the root, go down the tree and maintain:
  - Pointer  $x$ : traces the downward path (current node)
  - Pointer  $y$ : parent of  $x$  ("trailing pointer" )



# Example


Insert 13:





# Insert algorithm

```
1.  y ← NIL
2.  x ← root [T]
3.  while x ≠ NIL
4.    do y ← x
5.      if key [z] < key [x]
6.        x ← left [x]
7.      else
8.        x ← right [x]
9.  p[z] ← y
10. if y = NIL
11.   root [T] ← z    ▷ T was empty
12. else
13.   if key [z] < key [y]
14.     left [y] ← z
15.   else
16.     right [y] ← z
```



Best-case and worst-case time complexities?

Running time:  $O(h)$



# Exercise

---



- ▶ Build a binary search tree for the following sequence  
15, 6, 18, 3, 7, 17, 20, 2, 4





# Deletion

- ▶ Goal: Delete a given node  $z$  from a binary search tree
- ▶ Idea:
  - **Case 1:**  $z$  has no children
    - Delete  $z$  by making the parent of  $z$  point to NIL







# Deletion

## ► Case 2: z has one child

- Delete z by making the parent of z point to z's child, instead of to z, and link the parent with the new child






# Deletion

## ► Case 3: z has two children


- Find z's **successor** y (leftmost node in z's right subtree)
- y has either no or one right child (but no left child), why?
- Delete y from the tree (via Case 1 or 2)
- Replace z's key by y's key, and satellite data with y's





# Deletion algorithm


1. if  $\text{left}[z] = \text{NIL}$  and  $\text{right}[z] = \text{NIL}$  //z has no child
2.     if  $p[z] = \text{NIL}$  then  $\text{root}[T] = \text{NIL}$
3.     if  $z = \text{left}[p[z]]$
4.          $\text{left}[p[z]] = \text{NIL}$
5.     else
6.          $\text{right}[p[z]] = \text{NIL}$





# Deletion algorithm


```
1.  if left[z] = NIL and right[z] ≠ NIL  //z has one right child
2.      y = right[z]
3.      if p[z] = NIL
4.          root[T] = y
5.      else
6.          p[y] = p[z]
7.          if z = left[p[z]]
8.              left[p[z]] = y
9.          else
10.             right[p[z]] = y
11. if left[z] ≠ NIL and right[z] = NIL  //z has one left child
12.     y = left[z]
13.     if p[z] = NIL
14.         root[T] = y
15.     else
16.         p[y] = p[z]
17.         if z = left[p[z]]
18.             left[p[z]] = y
19.         else
20.             right[p[z]] = y
```





# Deletion algorithm

1. if  $\text{left}[z] \neq \text{NIL}$  and  $\text{right}[z] \neq \text{NIL}$  //z has two children
2.    $y \leftarrow \text{TREE-SUCCESSOR}(z)$    //left-most node in right tree
3.   if  $p[y] = z$
4.      $\text{right}[z] = \text{right}[y]$
5.     if  $\text{right}[y] \neq \text{NIL}$
6.        $p[\text{right}[y]] = z$
7.   else
8.     if  $\text{right}[y] = \text{NIL}$
9.        $\text{left}[p[y]] \leftarrow \text{NIL}$
10.    else
11.      $x \leftarrow \text{right}[y]$
12.      $p[x] \leftarrow p[y]$
13.      $\text{left}[p[y]] \leftarrow x$
14.    $\text{key}[z] \leftarrow \text{key}[y]$  //copy y's data into z



Best/worst-case time complexities?



# Summary

---

- ▶ Operations on binary search trees:
  - Search  $O(h)$
  - Predecessor  $O(h)$
  - Successor  $O(h)$
  - FindMin  $O(h)$
  - FindMax  $O(h)$
  - Insert/Delete  $O(h)$
  
- ▶ These operations are fast if the height of the tree is **small** - otherwise their performance is similar to that of a linked list



# Binary search trees vs linear lists

Operation	BST	Sorted-array-based List	Linked List
Constructor	$O(1)$	$O(1)$	$O(1)$
IsFull	$O(1)$	$O(1)$	$O(1)$
IsEmpty	$O(1)$	$O(1)$	$O(1)$
RetrieveItem	$O(\log N)^*$	$O(\log N)$	$O(N)$
InsertItem	$O(\log N)^*$	$O(N)$	$O(N)$
DeleteItem	$O(\log N)^*$	$O(N)$	$O(N)$

\*assuming  $h=O(\log N)$





# The issues in BST


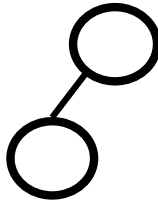
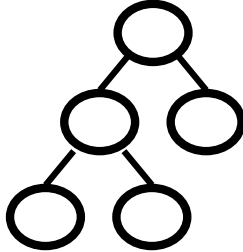
---

- ▶ After a series of delete operations, the above algorithm favors making **the left sub-trees deeper than the right**
- ▶ One solution:
  - Try to eliminate the problem by randomly choosing between the smallest element in the right sub-tree and the largest in the left when replacing the deleted element (not rigorous and not prove it yet!!)
- ▶ Existing balanced BST solutions
  - AVL tree: height  $O(\log n)$
  - Red-black tree: height  $O(\log n)$



# Exercise 1: count leaves

Example:

<p>A NULL binary tree has 0 leaf node</p>	 <p>A tree with 1 node has 1 leaf node</p>	 <p>No. of leaf nodes = 1</p>	 <p>No. of leaf nodes = 3</p>
---	---	--	--

//To count the number of leaf nodes


```
int Mytree::count_leaf(TreeNode* p)
{
    if (p == NULL)
        return 0;
    else if ((p->left == NULL) && (p->right == NULL))
        return 1;
    else
        return count_leaf(p->left) + count_leaf(p->right);
}
```



## Exercise 2: delete the node with two children

A bit complicated if we want to delete a NON-LEAF NODE with TWO children

1. Locate the node
2. Find the rightmost node in its left subtree
3. Or find the leftmost node in its right subtree
4. Use the key of the node to replace its key
5. Delete the node

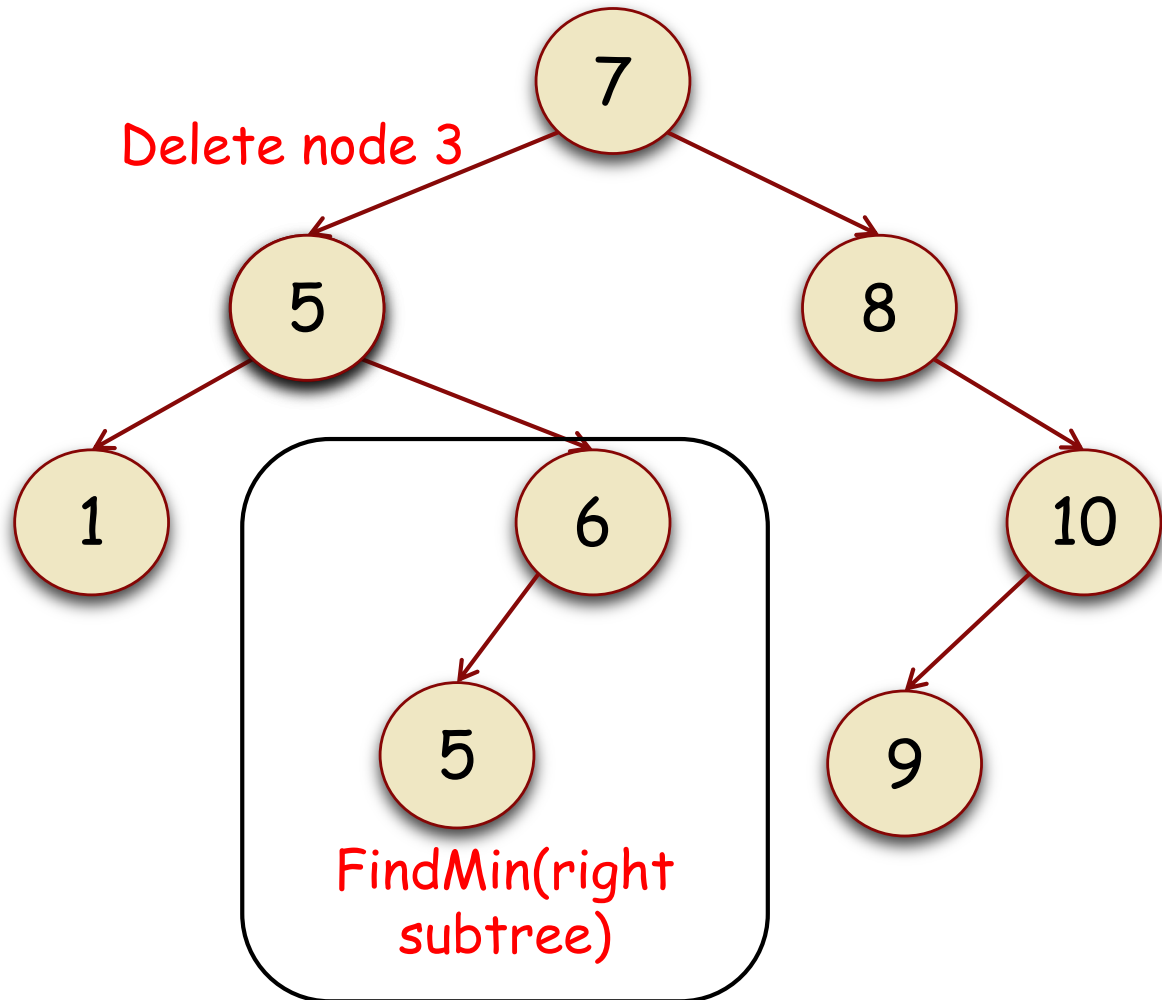




## Exercise 2: delete the node with two children

A bit complicated if we want to delete a NON-LEAF NODE with TWO children




1. Locate the node
2. Find the rightmost node in its left subtree
3. Or find the leftmost node in its right subtree
4. Use the key of the node to replace its key
5. Delete the node





## Exercise 3: operation commutative

- ▶ In a binary search tree, are the insert and delete operations commutative?
  - $\text{delete}(a)$  then  $\text{delete}(b) \Leftrightarrow \text{delete}(b)$  then  $\text{delete}(a)$ ?
  - $\text{insert}(a)$  then  $\text{insert}(b) \Leftrightarrow \text{insert}(b)$  then  $\text{insert}(a)$ ?



Case 1: Delete 5 and then 6  
Case 2: Delete 6 and then 5



## Exercise 4: sorting with BST

---

- ▶ How to sort an array of keys by building and traversing a BST?

```
1. Sort (A)
2.   for i = 1 to n
3.       insert(A[i])
4. inorder-tree-walk(root)
```


- What are the worst case and best case running times?
- In practice, how would this compare to other sorting algorithms?



# Exercise 5: lowest common ancestor

## ► Lowest common ancestor (LCA):

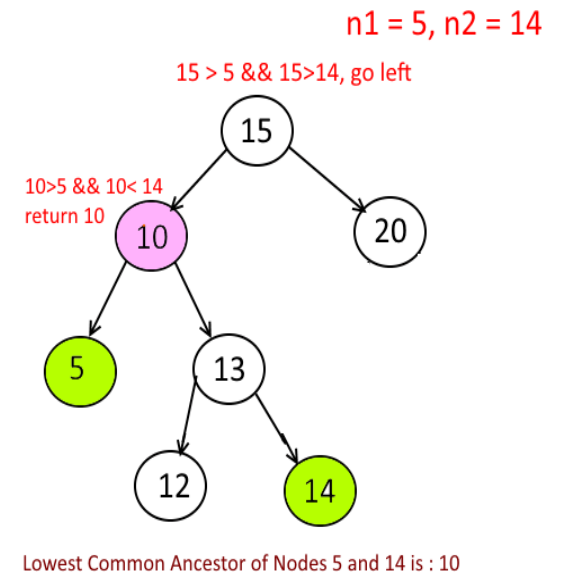
- The LCA of two nodes  $n1$  and  $n2$  is a node  $X$  such that node  $X$  will be the lowest node who has both  $n1$  and  $n2$  as its descendants
- Given a BST and two nodes  $n1$  and  $n2$ , how to find their LCA?



Lowest Ancestor Ancestor (5, 21) = 20  
Lowest Ancestor Ancestor (10, 30) = 25  
Lowest Ancestor Ancestor (5, 32) = 25  
Lowest Ancestor Ancestor (10, 23) = 20

Approach:

- 1) Start with the root
- 2) If  $root > n1$  and  $root > n2$  then lowest common ancestor will be in left subtree
- 3) If  $root < n1$  and  $root < n2$  then lowest common ancestor will be in right subtree
- 4) If Step 2 and Step 3 is false then we are at the root which is LCA, return it





# Recommended reading

---

- ▶ Reading this week
  - Chapter 12, textbook
- ▶ Next lecture
  - AVL-tree: Chapter 12, textbook