



香港中文大學 (深圳)
The Chinese University of Hong Kong

(Materials of this lecture are NOT included in the midterm and final exams)

CSC3100 Data Structures

Lecture 16: Red-black tree

Yixiang Fang
School of Data Science (SDS)
The Chinese University of Hong Kong, Shenzhen



Outline

- ▶ Definitions and examples
- ▶ Properties
- ▶ Operations
 - Insertion algorithm with three cases
 - Deletion algorithm (homework)



Red-black tree

- ▶ A “balanced” binary search tree
 - It guarantees an $O(\log n)$ running time for many operations, such as search, insertion, and deletion

- ▶ Overview
 - A binary search tree has an additional attribute for its nodes: color which can be either red or black
 - It restricts the way that nodes can be colored on any path from the root to a leaf
 - It ensures that no path is more than twice as long as any other path




Red-black tree properties

1. Every node is either red or black
2. The root is black
3. Every leaf (NIL) is black
4. If a node is red, then both its children are black
No two consecutive red nodes on a simple path from the root to a leaf
5. For each node, all paths from that node to descendant leaves contain the same number of black nodes




Example



- ▶ For convenience we use a sentinel $\text{NIL}[T]$ to represent all the NIL nodes at the leaves
 - $\text{NIL}[T]$ has the same fields as an ordinary node
 - $\text{Color}[\text{NIL}[T]] = \text{BLACK}$
 - The other fields may be set to arbitrary values



Black height of a node



► Height of a node x :

- $h(x)$ is the number of edges in the longest path to a leaf

► Black-height of a node x :

- $bh(x)$ is the number of black nodes (including NIL) on the path from x to a leaf, not counting x



Important property of red-black tree


A red-black tree with n internal nodes
has height at most $2\log(n + 1)$

- ▶ Need to prove two claims first ...



Claim 1


- ▶ Any node x with height $h(x)$ has $bh(x) \geq h(x)/2$
- ▶ **Proof**
 - By property 4, at most $h/2$ **red** nodes on the path from the node to a leaf
 - Hence at least $h/2$ are **black**





Claim 2

- ▶ The subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes





Claim 2 (Cont'd)

Proof: By induction on $h[x]$

Basis: $h[x] = 0 \Rightarrow$

x is a leaf ($NIL[T]$) \Rightarrow

$bh(x) = 0 \Rightarrow$

of internal nodes: $2^0 - 1 = 0$


Inductive hypothesis: assume it is true for $h[x] = h-1$



Claim 2 (Cont'd)

Inductive step:

- ▶ Prove it for $h[x] = h$
- ▶ Let $bh(x) = b$. Then, any child y of x has:
 - $bh(y) = b$ (if the child is **red**), or
 - $bh(y) = b - 1$ (if the child is **black**)





Claim 2 (Cont'd)


- ▶ Using inductive hypothesis, the number of internal nodes for each child of x is at least if x is black:

$$2^{bh(x) - 1} - 1$$

- ▶ What if x is red?

- ▶ The subtree rooted at x has at least:

$$\begin{aligned} & (2^{bh(x) - 1} - 1) + (2^{bh(x) - 1} - 1) + 1 \\ &= 2 \cdot (2^{bh(x) - 1} - 1) + 1 \\ &= 2^{bh(x)} - 1 \text{ internal nodes} \end{aligned}$$



$$bh(l) \geq bh(x) - 1$$

$$bh(r) \geq bh(x) - 1$$



Important property of red-black tree

A red-black tree with n internal nodes
has height at most $2\log(n + 1)$
Proof in the next slides.

- ▶ Claim 1: Any node x with height $h(x)$ has $bh(x) \geq h(x)/2$
- ▶ Claim 2: The subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes



Height of red-black tree


Lemma: A red-black tree with n internal nodes has height at most $2\log(n + 1)$.

Proof:

$$n \geq 2^{bh} - 1 \geq 2^{h/2} - 1$$

number n
of internal
nodes

since $bh \geq h/2$



► Add 1 to both sides and then take logs:

$$n + 1 \geq 2^{bh} \geq 2^{h/2}$$

$$\log(n + 1) \geq h/2$$

$$\Rightarrow h \leq 2 \log(n + 1)$$



Operations on red-black tree

- ▶ The non-modifying operations: **MINIMUM**, **MAXIMUM**, and **SEARCH** run in $O(h)$ time
 - They take $O(\log n)$ time on red-black trees
 - **SEARCH** is similar to the search on binary search tree
- ▶ What about **INSERT** and **DELETE**?
 - We have to guarantee that the modified tree will still be a red-black tree
 - Reconstruction will be too expensive
 - They can still be completed in $O(\log n)$ time



INSERT operation


INSERT: Suppose we want to insert 35. What color to make the new node?

► Red?

- Property 4 is violated: if a node is red, then its children are black

► Black?

- Property 5 is violated: all paths from a node to its leaves contain the same number of black nodes






Rotations

- ▶ After insertion and deletion on red-black trees, we need to restore the red-black tree properties
- ▶ Rotations take a red-black tree and a node within the tree and:
 - Two types of rotations: Left & right rotations
 - Together with some node re-coloring they help restore the red-black tree property
 - Change some of the pointer structure
 - Do not change the binary search tree property



Left and right rotations





INSERT

- ▶ Goal:
 - Insert a new node z into a red-black tree

- ▶ Idea:
 - Insert node z into the tree as for an ordinary BST

 - Color the node **red**

 - Restore the red-black-tree properties
 - Use an auxiliary procedure **RB-INSERT-FIXUP**




Properties affected by INSERT

1. Every node is either red or black OK!
2. The root is black If the root is changed
 \Rightarrow May not OK
3. Every leaf (NIL) is black OK!
4. If a node is red, then both its children are black

If $p(z)$ is red \Rightarrow not OK
 z and $p(z)$ are both red

- OK!
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes








INSERT(T, z)

1. $y \leftarrow \text{NIL}$
 2. $x \leftarrow \text{root}[T]$
 3. **while** $x \neq \text{NIL}$
 4. **do** $y \leftarrow x$
 5. **if** $\text{key}[z] < \text{key}[x]$
 6. **then** $x \leftarrow \text{left}[x]$
 7. **else** $x \leftarrow \text{right}[x]$
 8. $p[z] \leftarrow y$
 9. **if** $y = \text{NIL}$
 10. **then** $\text{root}[T] \leftarrow z$
 11. **else if** $\text{key}[z] < \text{key}[y]$
 12. **then** $\text{left}[y] \leftarrow z$
 13. **else** $\text{right}[y] \leftarrow z$
 14. $\text{left}[z] \leftarrow \text{NIL}$
 15. $\text{right}[z] \leftarrow \text{NIL}$
 16. $\text{color}[z] \leftarrow \text{RED}$
 17. $\text{RB-INSERT-FIXUP}(T, z)$
- Initialize nodes x and y
- Throughout the algorithm y points to the parent of x
- Go down the tree until reaching a leaf
- At that point y is the parent of the node to be inserted
- Sets the parent of z to be y
- The tree was empty: set the new node to be the root
- Otherwise, set z to be the left or right child of y , depending on whether the inserted node is smaller or larger than y 's key
- Set the fields of the newly added node
- Fix any inconsistencies that could have been introduced by adding this new red node



RB-Insert-Fixup(T, z)

- ▶ Case 1: z 's uncle y is red
 - Solution: recolor
- ▶ Case 2: z 's uncle y is black and z is a right child
 - Solution: double rotation
 - Can be transferred to Case 3
- ▶ Case 3: z 's uncle y is black and z is a left child
 - Solution: single rotation





RB-Insert-Fixup(T, z)

```
1. while z.p.color == red ← The while loop repeats only when
   Case 1 is executed: O(logn) times
2.     if z.p == z.p.p.left
3.         y = z.p.p.right
4.         if y.color == red
5.             z.p.color = black           // case 1
6.             y.color = black             // case 1
7.             z.p.p.color = red           // case 1
8.             z = z.p.p                   // case 1
9.         else if z == z.p.right
10.            z = z.p                      // case 2
11.            Left-rotation (T, z)         // case 2
12.            z.p.color = black            // case 3
13.            z.p.p.color = red            // case 3
14.            Right-rotation (T, z.p.p)    // case 3
15.     else (same as then clause with "right" and "left" exchanged)
16. T.root.color = black ← may just insert the root or the red violation reach root
```



INSERT: case 1




z's "uncle" (y) is **red**

Idea: (z is a right child)

- ▶ $p[p[z]]$ (z's grandparent) must be black: $p[z]$ is red

- ▶ Color $p[z]$ black
- ▶ Color y black
- ▶ Color $p[p[z]]$ **red**
- ▶ $z = p[p[z]]$

- Push the "**red**" violation up the tree





INSERT: case 1




z 's "uncle" (y) is **red**

Idea: (z is a left child)

- ▶ $p[p[z]]$ (z 's grandparent) must be black: $p[z]$ is red

- ▶ Color $p[z] \leftarrow$ black
- ▶ Color $y \leftarrow$ black
- ▶ Color $p[p[z]] \leftarrow$ **red**
- ▶ $z = p[p[z]]$

- Push the "**red**" violation up the tree








INSERT - case 3

Case 3:

- ▶ z's "uncle" (y) is black
- ▶ z is a left child

Idea:

- ▶ Color $p[z] \leftarrow \text{black}$
- ▶ Color $p[p[z]] \leftarrow \text{red}$
- ▶ $\text{RIGHT-ROTATE}(T, p[p[z]])$
 - ▶ No longer have 2 reds in a row
 - ▶ $p[z]$ is now black





INSERT - case 2



Case 2:

- ▶ z's "uncle" (y) is black
- ▶ z is a right child

Idea:


- ▶ $z \leftarrow p[z]$
- ▶ LEFT-ROTATE(T, z)

\Rightarrow now z is a left child, and both z and p[z] are red \Rightarrow case 3





Example





Complexity analysis

- ▶ Time complexity of detailed steps
 - A red-black tree has $O(\log n)$ height
 - Search for insertion location takes $O(\log n)$ time
 - Addition to the node takes $O(1)$ time
 - The while loop will be executed at most $O(\log n)$ time
 - Each recoloring and each rotation take $O(1)$ time
 - Never performs more than two rotations, since the loop terminates if case 2 or case 3 is executed
 - Hence, an insertion in a red-black tree takes $O(\log n)$ time

What are the advantages of red-black tree over AVL tree?



Exercises

- ▶ What is the ratio between the longest path and the shortest path in a red-black tree?
 - The shortest path is at least $bh(\text{root})$
 - The longest path is equal to $h(\text{root})$
 - Since $h(\text{root}) \leq 2bh(\text{root})$, the ratio is ≤ 2
- ▶ When we insert a node into a red-black tree, we initially set the color of the new node to red. Why didn't we choose to set the color to black?



Red-black trees: summary

- ▶ Red-black trees guarantee that the height of the tree will be $O(\log n)$

- ▶ Operations on red-black-trees:
 - SEARCH $O(h)$
 - PREDECESSOR $O(h)$
 - SUCCESSION $O(h)$
 - MINIMUM $O(h)$
 - MAXIMUM $O(h)$
 - INSERT $O(h)$
 - DELETE $O(h)$




Recommended reading

- ▶ Reading
 - Chapter 13, textbook
- ▶ Next lectures
 - Heap, chapters 6&12, textbook



DELETE operation




DELETE: the color of the node to be removed -- red

1. Every node is either red or black OK!
2. The root is black OK!
3. Every leaf (NIL) is black OK!
4. If a node is red, then both its children are black OK!
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes OK!

Note: the deletion of a red node is the same as the deletion of a node in BST



DELETE operation



DELETE: the color of the node to be removed -- **Black**

1. Every node is either **red** or **black** OK!
2. The root is **black** Not OK! If removing the root and the child that replaces it is **red**
3. Every leaf (NIL) is **black** OK!
4. If a node is **red**, then both its children are **black** Not OK! Could create two red nodes in a row

Not OK! Could change the black heights of some nodes

5. For each node, all paths from the node to descendant leaves contain the same number of **black** nodes




Deletion on red-black tree

- ▶ Similar to the deletion on BST, but need to use an auxiliary procedure **RB-Delete-Fixup** to restore the red-black tree properties
- ▶ Four different cases of **RB-Delete-Fixup**
 - Case 1: x's sibling w is red
 - Case 2: x's sibling w is black, and both of w's children are black
 - Case 3: x's sibling w is black, w's left child is red, and w's right child is black
 - Case 4: x's sibling w is black, and w's right child is red (left child either color)



Case 1


- ▶ Case 1: x's sibling w is red
 - Solution: rotate and recolor





Case 2


- ▶ Case 2: x 's sibling w is black, and both of w 's children are black
 - Solution: recolor





Case 3

- Case 3: x 's sibling w is black, w 's left child is red, and w 's right child is black





Case 4

- ▶ Case 4: x 's sibling w is black, and w 's right child is red (left child either color)

